

Gandalf vs Dumbledore

Introduction

In this challenge we will train a neural network to distinguish between two different types of images using convolutional neural networks.

In particular, we will train a model that can distinguish between images of Gandalf (from *Lord of the Rings*) and Dumbledore (from *Harry Potter*).

This is a challenging visual task for a few reasons:

- Both are grey-bearded old white men
- Both wear wizard robes and hats
- Two different actors played Dumbledore (yes, we will ignore the Jude Law incarnation), Richard Harris and Michael Gambon, and they did not look very much alike. In fact, you could argue that Michael Gambon looks much more like Ian McKellen than like Richard Harris.
- Gandalf has two incarnations that look significantly different: Gandalf the Grey and Gandalf the White

With a very limited training set of images, this task would be close to impossible if we were training a neural network from scratch.

Thankfully, but we will not start from scratch. We will use Transfer Learning to benefit from the representations learned by other networks previously trained on huge datasets with many image classes, and apply those representations to our much more specific problem, with much more limited data.

Suggested resources

- Easy way to manipulate an image dataset with Keras: <https://keras.io/api/preprocessing/image/>
- You will probably find this Transfer Learning example in Keras very useful: https://keras.io/examples/vision/image_classification_efficientnet_fine_tuning/
- Data augmentation layers in Keras: https://www.tensorflow.org/api_docs/python/tf/keras/layers/experimental/preprocessing
- There are many pre-trained state-of-art models available directly through the Keras API, such as EfficientNet: <https://keras.io/api/applications/efficientnet/>

Setup

First, we will install and load all of the libraries we will be using.

Please update it as you need, importing whatever libraries you use in your code.

1. Data Analysis

Include any imports you need in the cell below.

I have included some imports to give you hints on what you might want to make use of.

```

1 import tensorflow as tf
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 from PIL import Image
6 from google_drive_downloader import GoogleDriveDownloader as gdd
7 from sklearn.model_selection import train_test_split
8 import shutil
9 from tensorflow import keras
10 from tensorflow.keras import layers
11 import pandas as pd
12 from sklearn import preprocessing
13 from keras.utils import np_utils
14 from keras.models import Sequential
15 from keras.layers import Dense, Dropout
16 from keras.layers import Activation
17 from keras.callbacks import EarlyStopping
18 SEED = 1534 # Set this to whichever value you want and use this parameters in any seeded pseudorandom

```

Dataset

You can download the training data by running the cell below. You will get the following folder structure:

```

train_data
|_ dumbledore
    |_ dumbledore_0000.jpg
    |_ dumbledore_0001.jpg
    |_ dumbledore_0002.jpg
    |_ ...
|_ gandalf
    |_ gandalf_0000.jpg
    |_ gandalf_0001.jpg
    |_ gandalf_0002.jpg

```

► Load the data: Dumbledore and Gandalf dataset

```
[ ] ↳ 1 celda oculta
```

▼ Filter out corrupted images

```

1 import os
2
3 num_skipped = 0
4 for folder_name in ("dumbledore", "gandalf"):
5     folder_path = os.path.join("/content/drive/MyDrive/datasets/DL2_GvD/train_data", folder_name)
6     for fname in os.listdir(folder_path):
7         fpath = os.path.join(folder_path, fname)
8         try:
9             fobj = open(fpath, "rb")

```

```

10         is_jfif = tf.compat.as_bytes("JFIF") in fobj.peek(10)
11     finally:
12         fobj.close()
13
14     if not is_jfif:
15         num_skipped += 1
16         # Delete corrupted image
17         os.remove(fpath)
18
19 print("Deleted %d images" % num_skipped)

```

Deleted 0 images

▼ Generate a Dataset

```

1 image_size = (180, 180)
2 batch_size = 32 # El bach no puede ser mucho mayor debido al tamaño del dataset
3 validation_split = 0.2 # 20%del dataset para validacion
4
5 # Se crea el conjunto de datos
6 train_ds = tf.keras.preprocessing.image_dataset_from_directory(
7     "train_data",
8     validation_split=0.2,
9     subset="training",
10    seed=1337,
11    image_size=image_size,
12    batch_size=batch_size,
13 )
14
15 # Se crea el conjunto de validacion
16 val_ds = tf.keras.preprocessing.image_dataset_from_directory(
17     "train_data",
18     validation_split=0.2,
19     subset="validation",
20     seed=1337,
21     image_size=image_size,
22     batch_size=batch_size,
23 )
24 train_ds

```

Found 80 files belonging to 2 classes.

Using 64 files for training.

Found 80 files belonging to 2 classes.

Using 16 files for validation.

<BatchDataset element_spec=(TensorSpec(shape=(None, 180, 180, 3), dtype=tf.float32, name=None))

▼ Visualize the data

```

1 import matplotlib.pyplot as plt
2
3 plt.figure(figsize=(10, 10))
4 for images, labels in train_ds.take(1):
5     for i in range(9):

```

```

6     ax = plt.subplot(3, 3, i + 1)
7     plt.imshow(images[i].numpy().astype("uint8"))
8     plt.title(int(labels[i]))
9     plt.axis("off")

```



▼ 2. Data Processing

▼ Using image data augmentation

When you don't have a large image dataset, it's a good practice to artificially introduce sample diversity by applying random yet realistic transformations to the training images, such as random horizontal flipping or small random rotations. This helps expose the model to different aspects of the training data while slowing down overfitting.

```

1 image_augmentation = keras.Sequential(
2     [
3         layers.RandomFlip(),
4         layers.RandomRotation(factor=0.1),
5         layers.RandomContrast(0.01),
6         #layers.RandomTranslation(height_factor=0.1, width_factor=0.1, seed=SEED),
7         #layers.RandomZoom(0.2, fill_mode='nearest', seed=SEED),
8     ],

```

```
9     name = "image_augmentation"  
10 )
```

Let's visualize what the augmented samples look like, by applying `data_augmentation` repeatedly to the first image in the dataset:

```
1 plt.figure(figsize=(10, 10))  
2 for images, _ in train_ds.take(1):  
3     for i in range(9):  
4         augmented_images = image_augmentation(images)  
5         ax = plt.subplot(3, 3, i + 1)  
6         plt.imshow(augmented_images[0].numpy().astype("uint8"))  
7         plt.axis("off")
```



▼ 3. Build a model

▼ Model

I have chosen to use the ResNet152 architecture, replacing the top layers to adjust it to our binary classification problem. We have the choice to import either a blank model or a pretrained one, depending on the train parameter. We will shortly see the difference of performance between both options.

```
1 # ResNet152
2 from keras.applications.resnet import ResNet152 as Architecture
3
4 def build_model(input_shape, num_classes, trained):# Depending on the train parameter, it returns
5     base_model = Architecture (include_top=False, input_shape=input_shape, weights='imagenet'if tra
6     base_model.trainable = not trained
7
8     inputs = keras.Input(shape=input_shape, name='input')
9     x = image_augmentation(inputs)
10    x = base_model(x)
11    x = layers.GlobalAveragePooling2D(name='avg_pool')(x)
12    x = layers.BatchNormalization()(x)
13    x = layers.Dropout(0.2,name='top_dropout')(x)
14
15    if num_classes == 2:
16        activation = 'sigmoid'
17        units = 1
18    else:
19        activation = 'softmax'
20        units = num_classes
21
22    outputs = layers.Dense(units,activation=activation, name='pred')(x)
23    return keras.Model(inputs, outputs, name=Architecture.__name__)
```

```
1 # function to plot the accuracy
2 def plot_model_hist(model_history):
3     sns.set()
4     fig, axs = plt.subplots(1,2, figsize=(15,5))
5     # Summarize history for accuracy
6     axs[0].plot(range(1,len(model_history.history['accuracy'])+1), model_history.history['accuracy'])
7     axs[0].plot(range(1,len(model_history.history['val_accuracy'])+1), model_history.history['val_accuracy'])
8     axs[0].set_title('Model Accuracy')
9     axs[0].set_ylabel('Accuracy')
10    axs[0].set_xlabel('Epoch')
11    axs[0].legend(['train','val'],loc='best')
12
13    # Summarize history for loss
14    axs[1].plot(range(1,len(model_history.history['loss'])+1), model_history.history['loss'])
15    axs[1].plot(range(1,len(model_history.history['val_loss'])+1), model_history.history['val_loss'])
16    axs[1].set_title('Model Loss')
17    axs[1].set_ylabel('Loss')
18    axs[1].set_xlabel('Epoch')
19    axs[1].legend(['train','val'],loc='best')
20
21    plt.show()
```

▼ Blanck Model

In this model, the ResNet 152 presents random weights, that is, it does not have pre-trained parameters.

Total parameters: 58,381,185

Trainable parameters: 58,225,665

```
1 model = build_model(image_size+(3,), 2, trained=False)
2 model.summary()
```

Model: "ResNet152"

Layer (type)	Output Shape	Param #
=====		
input (InputLayer)	[(None, 180, 180, 3)]	0
image_augmentation (Sequential)	(None, 180, 180, 3)	0
resnet152 (Functional)	(None, 6, 6, 2048)	58370944
avg_pool (GlobalAveragePooling2D)	(None, 2048)	0
batch_normalization_2 (Batch Normalization)	(None, 2048)	8192
top_dropout (Dropout)	(None, 2048)	0
pred (Dense)	(None, 1)	2049
=====		
Total params: 58,381,185		
Trainable params: 58,225,665		
Non-trainable params: 155,520		
=====		

```
1 epochs = 50
2
3 callbacks = [
4     keras.callbacks.EarlyStopping(patience=10, restore_best_weights=True),
5     keras.callbacks.ModelCheckpoint(save_best_only=True, mode='min', filepath='model.hdf5')
6 ]
7
8 model.compile(
9     optimizer=keras.optimizers.Adam(1e-2),
10    loss='binary_crossentropy',
11    metrics=['accuracy']
12 )
13 history=model.fit(train_ds, epochs=epochs, callbacks=callbacks, validation_data=val_ds)
14 plot_model_hist(history)
```



```

Epoch 1/50
2/2 [=====] - 127s 54s/step - loss: 4.4609 - accuracy: 0.7188 - v
Epoch 2/50
2/2 [=====] - 94s 49s/step - loss: 1.6827 - accuracy: 0.5469 - va
Epoch 3/50
2/2 [=====] - 94s 49s/step - loss: 0.8310 - accuracy: 0.5625 - va
Epoch 4/50
2/2 [=====] - 93s 49s/step - loss: 0.7864 - accuracy: 0.5156 - va
Epoch 5/50
2/2 [=====] - 92s 48s/step - loss: 0.9432 - accuracy: 0.5000 - va
Epoch 6/50
2/2 [=====] - 98s 54s/step - loss: 1.6198 - accuracy: 0.4844 - va
Epoch 7/50
2/2 [=====] - 95s 49s/step - loss: 1.4286 - accuracy: 0.5312 - va
Epoch 8/50
2/2 [=====] - 93s 49s/step - loss: 1.8326 - accuracy: 0.5000 - va
Epoch 9/50
2/2 [=====] - 93s 49s/step - loss: 2.7686 - accuracy: 0.4844 - va
Epoch 10/50
2/2 [=====] - 93s 49s/step - loss: 0.7260 - accuracy: 0.5312 - va

```



As we can see, we do not seem to have enough images to train a model a from scratch, resulting in a random model with 50% accuracy on the validation set. Lets see how a pretrained model performs instead



▼ Pretrained Model

```

1 model = build_model(image_size+(3,), 2,trained=True)
2 model.summary()

```

Model: "ResNet152"

Layer (type)	Output Shape	Param #
=====		
input (InputLayer)	[(None, 180, 180, 3)]	0
image_augmentation (Sequential)	(None, 180, 180, 3)	0
resnet152 (Functional)	(None, 6, 6, 2048)	58370944
avg_pool (GlobalAveragePooling2D)	(None, 2048)	0
batch_normalization_3 (Batch Normalization)	(None, 2048)	8192
top_dropout (Dropout)	(None, 2048)	0
pred (Dense)	(None, 1)	2049
=====		

Total params: 58,381,185
Trainable params: 6,145
Non-trainable params: 58,375,040

```
1 epochs = 50
2
3 callbacks = [
4     keras.callbacks.EarlyStopping(patience=20, restore_best_weights=True),
5     keras.callbacks.ModelCheckpoint(save_best_only=True, mode='min', filepath='petrained_model.hd
6 ]
7
8 model.compile(
9     optimizer=keras.optimizers.Adam(2e-3),
10    loss='binary_crossentropy',
11    metrics=['accuracy']
12 )
13 history=model.fit(train_ds, epochs=epochs, callbacks=callbacks, validation_data=val_ds)
14 plot_model_hist(history)
```

Epoch 1/50
2/2 [=====] - 42s 20s/step - loss: 1.0535 - accuracy: 0.4375 - va
Epoch 2/50
2/2 [=====] - 27s 16s/step - loss: 0.8007 - accuracy: 0.5469 - va
Epoch 3/50
2/2 [=====] - 27s 16s/step - loss: 0.5551 - accuracy: 0.7500 - va
Epoch 4/50
2/2 [=====] - 26s 16s/step - loss: 0.3965 - accuracy: 0.8594 - va
Epoch 5/50
2/2 [=====] - 27s 16s/step - loss: 0.2998 - accuracy: 0.8438 - va
Epoch 6/50
2/2 [=====] - 27s 16s/step - loss: 0.3973 - accuracy: 0.8438 - va
Epoch 7/50
2/2 [=====] - 27s 17s/step - loss: 0.2735 - accuracy: 0.8750 - va
Epoch 8/50
2/2 [=====] - 28s 17s/step - loss: 0.3505 - accuracy: 0.8281 - va
Epoch 9/50
2/2 [=====] - 25s 15s/step - loss: 0.2062 - accuracy: 0.8906 - va
Epoch 10/50
2/2 [=====] - 27s 16s/step - loss: 0.2336 - accuracy: 0.9219 - va
Epoch 11/50
2/2 [=====] - 27s 17s/step - loss: 0.1981 - accuracy: 0.9062 - va
Epoch 12/50
2/2 [=====] - 27s 16s/step - loss: 0.1679 - accuracy: 0.9375 - va
Epoch 13/50
2/2 [=====] - 27s 16s/step - loss: 0.1447 - accuracy: 0.9531 - va
Epoch 14/50
2/2 [=====] - 26s 16s/step - loss: 0.1432 - accuracy: 0.9531 - va
Epoch 15/50
2/2 [=====] - 27s 17s/step - loss: 0.0910 - accuracy: 0.9844 - va
Epoch 16/50
2/2 [=====] - 32s 16s/step - loss: 0.1264 - accuracy: 0.9531 - va
Epoch 17/50
2/2 [=====] - 27s 16s/step - loss: 0.0766 - accuracy: 0.9844 - va
Epoch 18/50
2/2 [=====] - 26s 16s/step - loss: 0.0503 - accuracy: 1.0000 - va
Epoch 19/50
2/2 [=====] - 26s 16s/step - loss: 0.0880 - accuracy: 0.9844 - va
Epoch 20/50
2/2 [=====] - 27s 16s/step - loss: 0.0999 - accuracy: 0.9844 - va
Epoch 21/50
2/2 [=====] - 27s 17s/step - loss: 0.0690 - accuracy: 0.9844 - va
Epoch 22/50
2/2 [=====] - 27s 16s/step - loss: 0.0934 - accuracy: 0.9688 - va
Epoch 23/50
2/2 [=====] - 26s 16s/step - loss: 0.0589 - accuracy: 1.0000 - va
Epoch 24/50
2/2 [=====] - 25s 15s/step - loss: 0.0355 - accuracy: 1.0000 - va
Epoch 25/50
2/2 [=====] - 25s 15s/step - loss: 0.0666 - accuracy: 0.9688 - va
Epoch 26/50
2/2 [=====] - 25s 15s/step - loss: 0.0294 - accuracy: 1.0000 - va
Epoch 27/50
2/2 [=====] - 25s 15s/step - loss: 0.0514 - accuracy: 1.0000 - va
Epoch 28/50
2/2 [=====] - 25s 15s/step - loss: 0.0409 - accuracy: 1.0000 - va
Epoch 29/50
2/2 [=====] - 25s 15s/step - loss: 0.0732 - accuracy: 0.9844 - va
Epoch 30/50
2/2 [=====] - 25s 15s/step - loss: 0.0735 - accuracy: 0.9844 - va
Epoch 31/50
2/2 [=====] - 25s 15s/step - loss: 0.0582 - accuracy: 0.9844 - va
Epoch 32/50
2/2 [=====] - 25s 15s/step - loss: 0.0515 - accuracy: 0.9844 - va

```

Epoch 33/50
2/2 [=====] - 25s 15s/step - loss: 0.0343 - accuracy: 0.9844 - va
Epoch 34/50
2/2 [=====] - 25s 15s/step - loss: 0.0387 - accuracy: 1.0000 - va
Epoch 35/50
2/2 [=====] - 25s 15s/step - loss: 0.0324 - accuracy: 1.0000 - va
Epoch 36/50
2/2 [=====] - 25s 15s/step - loss: 0.0668 - accuracy: 0.9688 - va
Epoch 37/50
2/2 [=====] - 25s 15s/step - loss: 0.0285 - accuracy: 1.0000 - va
Epoch 38/50
2/2 [=====] - 25s 15s/step - loss: 0.0248 - accuracy: 1.0000 - va
Epoch 39/50
2/2 [=====] - 25s 15s/step - loss: 0.0466 - accuracy: 1.0000 - va
Epoch 40/50
2/2 [=====] - 27s 17s/step - loss: 0.0388 - accuracy: 1.0000 - va
Epoch 41/50
2/2 [=====] - 27s 16s/step - loss: 0.0321 - accuracy: 1.0000 - va
Epoch 42/50

```

As we can observe from these results, using a pretrained model allows us to focus on the training the last layer, used for classification while letting to bottom layers frozen. The bottom layers seem to be extracting good features for our top layers to use, but could be do better?

Fine Turing

```

2/2 [=====] - 27s 16s/step - loss: 0.0349 - accuracy: 1.0000 - va

```

We unfreeze the bottom layers of our model and train them, hopefully improving the predicting power of our model

```

Epoch 43/50
2/2 [=====] - 27s 16s/step - loss: 0.0349 - accuracy: 1.0000 - va

```

```

1 def unfreeze_model(model):
2 # We unfreeze the top 20 layers while leaving BatchNorm layers frozen
3 for layer in model.layers[-20:]:
4     if not isinstance(layer, layers.BatchNormization):
5         layer.trainable=True
6 return model

```

```

1 model = unfreeze_model(model)
2 model.summary()

```


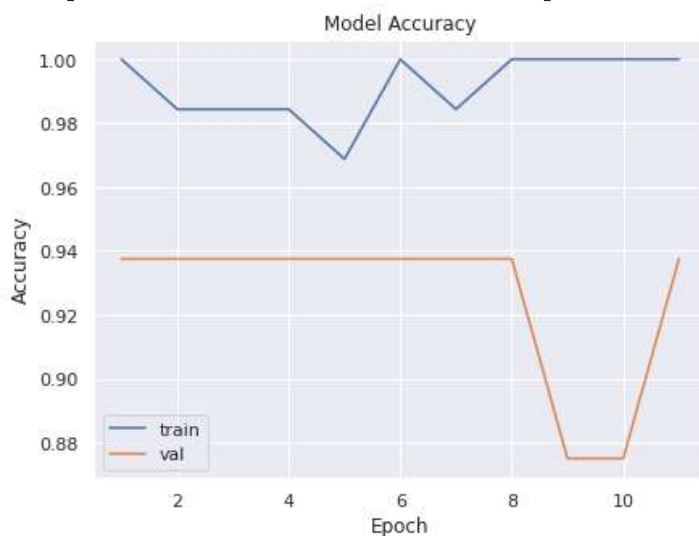
Model: "ResNet152"

Layer (type)	Output Shape	Param #
input (InputLayer)	[(None, 180, 180, 3)]	0
image_augmentation (Sequent ial)	(None, 180, 180, 3)	0
resnet152 (Functional)	(None, 6, 6, 2048)	58370944
avg_pool (GlobalAveragePool ing2D)	(None, 2048)	0
batch_normalization_3 (Batc hNormalization)	(None, 2048)	8192
top_dropout (Dropout)	(None, 2048)	0

```
=====
Total params: 58,381,185
Trainable params: 58,225,665
Non-trainable params: 155,520
=====
```

```
1 epochs = 50
2
3 callbacks = [
4     keras.callbacks.EarlyStopping(patience=10, restore_best_weights=True),
5     keras.callbacks.ModelCheckpoint(save_best_only=True, mode='min', filepath='finetunedpettrained
6 ]
7 history=model.fit(train_ds, epochs=epochs, callbacks=callbacks, validation_data=val_ds)
8 plot_model_hist(history)
```

```
Epoch 1/50
2/2 [=====] - 28s 18s/step - loss: 0.0387 - accuracy: 1.0000 - val_loss: 0.2050
Epoch 2/50
2/2 [=====] - 25s 15s/step - loss: 0.0494 - accuracy: 0.9844 - val_loss: 0.2100
Epoch 3/50
2/2 [=====] - 25s 15s/step - loss: 0.0586 - accuracy: 0.9844 - val_loss: 0.2150
Epoch 4/50
2/2 [=====] - 25s 15s/step - loss: 0.0599 - accuracy: 0.9844 - val_loss: 0.2200
Epoch 5/50
2/2 [=====] - 25s 15s/step - loss: 0.0456 - accuracy: 0.9688 - val_loss: 0.2250
Epoch 6/50
2/2 [=====] - 25s 15s/step - loss: 0.0363 - accuracy: 1.0000 - val_loss: 0.2300
Epoch 7/50
2/2 [=====] - 25s 15s/step - loss: 0.0432 - accuracy: 0.9844 - val_loss: 0.2350
Epoch 8/50
2/2 [=====] - 25s 15s/step - loss: 0.0272 - accuracy: 1.0000 - val_loss: 0.2400
Epoch 9/50
2/2 [=====] - 25s 15s/step - loss: 0.0175 - accuracy: 1.0000 - val_loss: 0.2450
Epoch 10/50
2/2 [=====] - 25s 15s/step - loss: 0.0177 - accuracy: 1.0000 - val_loss: 0.2500
Epoch 11/50
2/2 [=====] - 26s 15s/step - loss: 0.0298 - accuracy: 1.0000 - val_loss: 0.2550
```



Haz doble clic (o pulsa Intro) para editar

1